

# MEC302: Embedded Computer Systems

## Theme II: Design of Embedded Computer Systems

### Lecture 8 – Multitasking

Dr. Timur Saifutdinov

Assistant Professor at EEE, SAT

Email:

Timur.Saifutdinov@xjtlu.edu.cn

# Outline

- Multitasking
  - What is multitasking;
  - Multitasking mechanisms;
- Threads
  - Creating and handling threads;
  - Memory consistency models of threads;
  - Mutual exclusions
- Processes
  - Creating processes;
  - Communication between processes;
  - Example of files handling

# What is multitasking

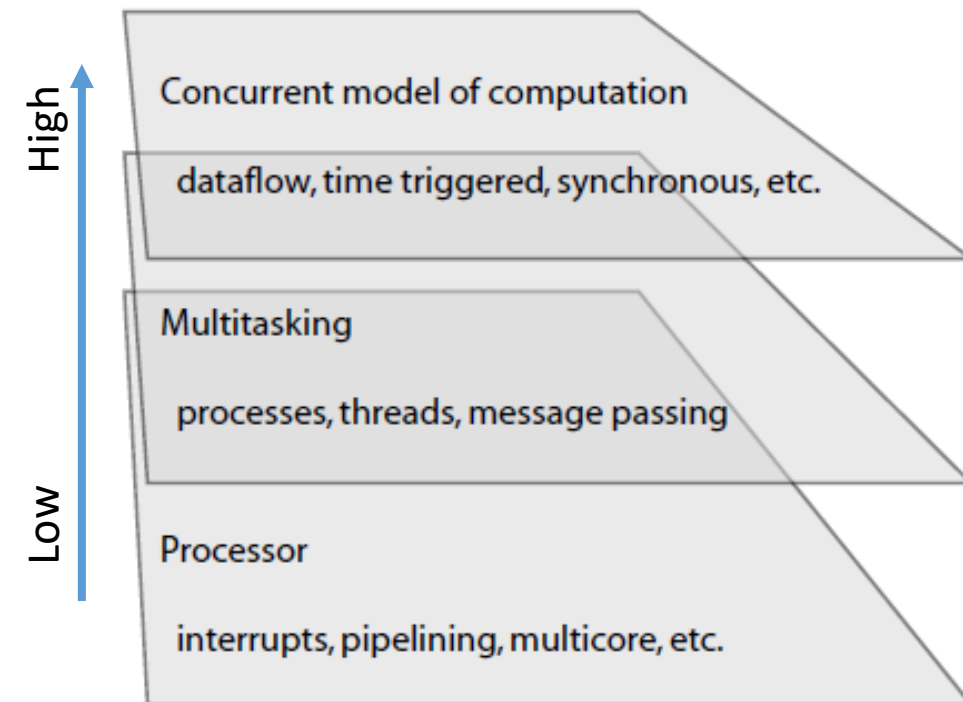
**Multitasking** – software mechanisms that provides concurrent execution of sequential code (i.e., simultaneous execution of multiple tasks).

Reasons for concurrency (i.e., **advantages**):

- Increase responsiveness (reduce latency) to external signal (stimuli);
- Improve performance when running on multiple processors or cores;
- Control timing of external interactions (e.g., update display while executing other tasks).

**Multitasking** bridges the high-level of abstract computational and dataflow models and low-level hardware mechanisms.

Layers of abstraction for concurrency in programs:



# Multitasking mechanisms


**Multitasking** is provided by an **Operating System (OS)** through:

- Library of procedures;
- (Micro)**kernel**.

**OS** is a software that manages computer hardware and software resources to provide interface between user and computer.

**OS** can differ by application:

- **General-purpose**: MS Windows, Mac OS X, Linux;
- **Real-time OS (RTOS)** for embedded applications: WinCE, QNX, UNIX-RT;
- **Mobile OS** for handheld devices: Symbian OS, iOS, Android.



The core of an OS is **kernel**, which controls all programs running on the computer:

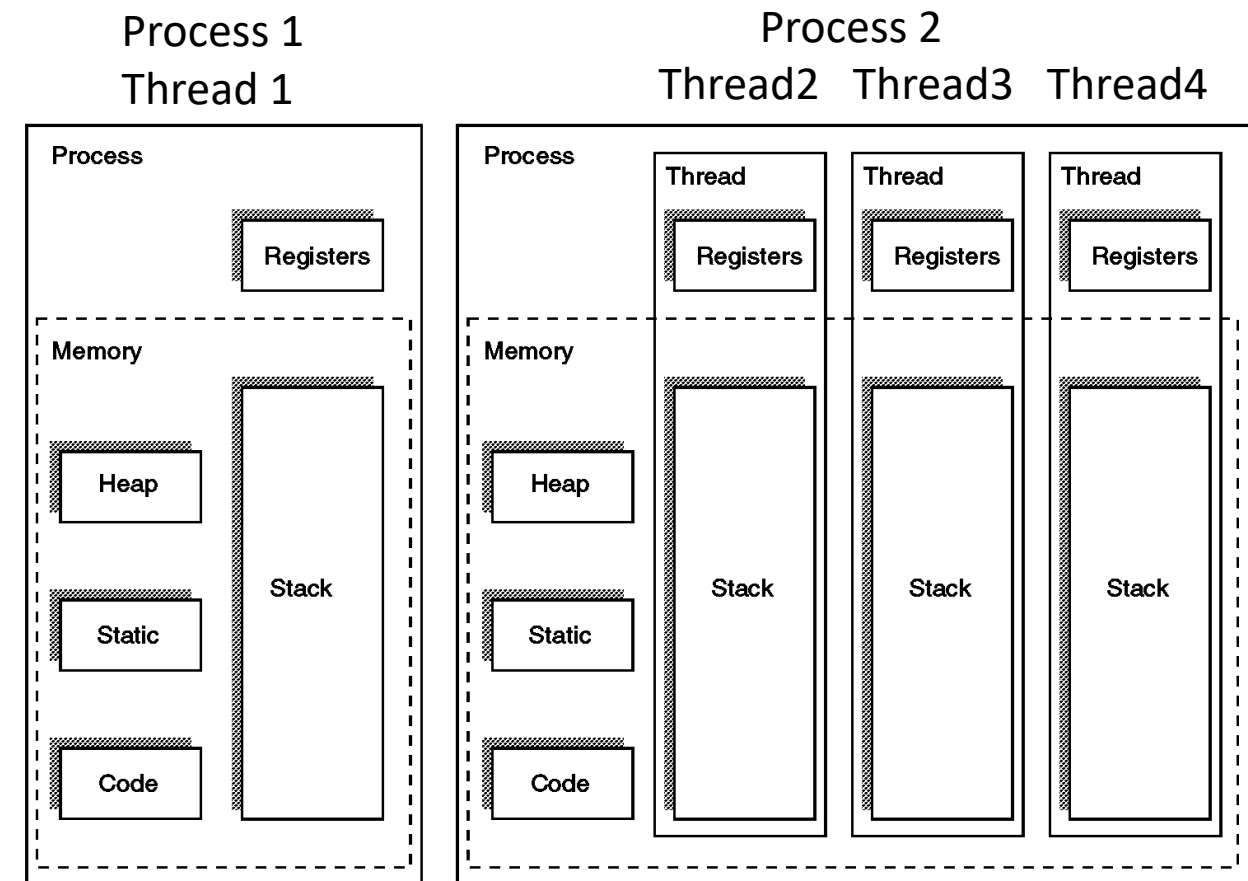
- Converts user commands to machine instructions;
- Allocate memory for processes;
- Communicate with peripherals and networks;
- Decide on the order of processes execution (i.e., **scheduler** for multiple tasks).

# Multitasking mechanisms

**Multitasking mechanisms** provided by an **OS** include:

- **Threads** are imperative programs that run concurrently and share a memory space. They can:
  - Directly access each others' variables.
- **Processes** are imperative programs with their own memory spaces. Communication is done via:
  - File system – exchange data through files;
  - Message passing – controlled blocks of memory

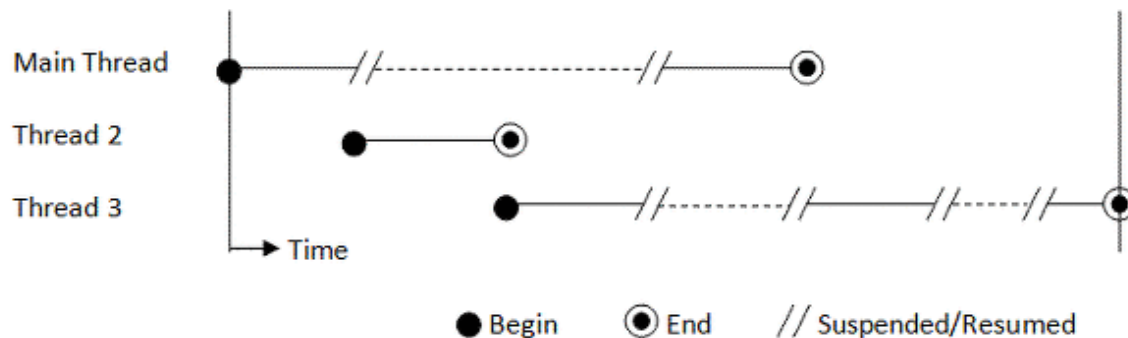
#Memory allocation for processes and threads [2]:



# Handling threads and processes (i.e., tasks)

**Scheduler** is a procedure (provided by a **kernel**) that decides which task(s) to execute next, e.g., based on:

- **Fairness** – equal opportunity for a task routines to be called;
- **Timing** – amount of time given for a task;
- **Priority** – a lower priority task is executed only after all higher priority tasks terminated;
- Other methods (to be studied **next lecture**)



**Techniques** to invoke the **scheduler** procedure can follow:

- **Cooperative multitasking** – does not interrupt a task unless the thread/process itself calls a certain procedure or terminates;
- **Interrupt Service Routine (ISR)** – tasks are interrupted by the timer every fixed duration of time (i.e., **jiffy**):
  - Small **jiffy** (eg,  $1\mu\text{s}$ ) can degrade overall performance;
  - Large **jiffy** (eg,  $10\text{ms}$ ) compromise real-time response.

# Creating threads

Thread library supports using **threads**, e.g.:

- **POSIX threads (Pthreads)** is standardized Application Program Interface (API) supported by many OS (including WinCE, Unix).

To create a **thread** you need:

- Define processor directive:
  - *#include <pthread.h>*
- Declare ID, create and exit status variables:
  - *pthread\_t ID;*
  - *int createStatus;*
  - *void\* exitStatus;*
- Declare and specify thread procedure (i.e., function):
  - *void fun\_name(void){...}*
- Create thread:
  - *createStatus = pthread\_create(&ID,&attr,fun\_name,&arg);*

#Consider a multi-thread C program:

```
1  #include <pthread.h>
2  #include <stdio.h>
3  void* printN(void* arg) { // Procedure assigned to threads
4      int i;
5      for (i = 0; i < 10; i++) {
6          printf("My ID: %d\n", *(int*)arg);
7      }
8      return NULL;
9  }
10 int main(void) {
11     pthread_t threadID1, threadID2;
12     void* exitStatus;
13     int x1 = 1, x2 = 2;
14     pthread_create(&threadID1, NULL, printN, &x1);
15     pthread_create(&threadID2, NULL, printN, &x2);
16     printf("Started threads.\n");
17     pthread_join(threadID1, &exitStatus);
18     pthread_join(threadID2, &exitStatus);
19     return 0;
20 }
```

# Memory consistency models of threads

All possible implementation of threads\* defines a **Memory Consistency (MC) model**.

**MC model** defines how variables that are read and written by different threads can be accessed by those threads (e.g., allowed ordering of access).

\* - A thread may be suspended between any two **atomic operations** to execute another thread and/or an **ISR**.

#Consider example:

1. w, x, y, z = 0;

Thread A:

☐ x = 1;

☐ w = y;

Thread B:

☐ y = 1;

☐ z = x;

Depending on the execution order, we:

- Know that x, y and {w or z} are ones – **sequential consistency**;
- If the order of threads execution affects program results, then it is called **Race conditions**.

**!Race conditions** can be disastrous for applications, but not necessarily. 8



# Mutual exclusion (mutex)

A **mutual exclusion lock (mutex)** prevents any two threads from simultaneously accessing or modifying a shared resource (i.e., memory).

**Mutexes** are initialized by creating an instance of a structure:

- `pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;`

The lock is acquired by a thread through:

- `pthread_mutex_lock(&lock);`

The lock is released by a thread through:

- `pthread_mutex_unlock(&lock);`

**Deadlocks** (side-effect of **mutex**) may occur when some threads become permanently blocked trying to acquire locks held by other threads:

- May occur when using multiple locks.

#Consider the same example with mutex:

1. `w, x, y, z = 0;`
2. `pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;`

Thread A:

- ☐ `pthread_mutex_lock(&lock);`
- ☐ `x = 1;`
- ☐ `w = y;`
- ☐ `pthread_mutex_unlock(&lock);`

Thread B:

- ☐ `pthread_mutex_lock(&lock);`
- ☐ `y = 1;`
- ☐ `z = x;`
- ☐ `pthread_mutex_unlock(&lock);`

#Consider the example with multiple locks:

1. `w, x, y, z = 0;`
2. `pthread_mutex_t lock1 = PTHREAD_MUTEX_INITIALIZER;`
3. `pthread_mutex_t lock2 = PTHREAD_MUTEX_INITIALIZER;`

Thread A:

- ☐ `pthread_mutex_lock(&lock1);`
- ☐ `x = 1;`
- ☐ `pthread_mutex_lock(&lock2);`
- ☐ `w = y;`
- ☐ `pthread_mutex_unlock(&lock1);`
- ☐ `pthread_mutex_unlock(&lock2);`

Thread B:

- ☐ `pthread_mutex_lock(&lock2);`
- ☐ `y = 1;`
- ☐ `pthread_mutex_lock(&lock1);`
- ☐ `z = x;`
- ☐ `pthread_mutex_unlock(&lock2);`
- ☐ `pthread_mutex_unlock(&lock1);`

# Mutual exclusion (mutex)

## Techniques to avoid **deadlocks**:

- **Use only one lock** – may not be practical for concurrency (i.e. real-time constraints and program modularity);
- **Avoid nested locks** – equivalent to the above;
- **Use global mutex** to disable blocks at some frequency to pass deadlocks;
- **Use a hierarchy of locks** – acquire locks in the same order in all threads that use them.
- Other methods...

#Consider example with the hierarchy of locks:

1. `w, x, y, z = 0;`
2. `pthread_mutex_t lock1 = PTHREAD_MUTEX_INITIALIZER;`
3. `pthread_mutex_t lock2 = PTHREAD_MUTEX_INITIALIZER;`

Thread A:

```
☐ pthread_mutex_lock(&lock1);  
☐ x = 1;  
☐ pthread_mutex_lock(&lock2);  
☐ w = y;  
☐ pthread_mutex_unlock(&lock2);  
☐ pthread_mutex_unlock(&lock1);
```

Thread B:

```
☐ pthread_mutex_lock(&lock1);  
☐ y = 1;  
☐ pthread_mutex_lock(&lock2);  
☐ z = x;  
☐ pthread_mutex_unlock(&lock2);  
☐ pthread_mutex_unlock(&lock1);
```

# Creating processes

**Processes** are imperative programs with their own memory spaces that can run concurrently.

#In UNIX, two or more programs can be run from the command line (i.e., bash) as follows:

*prog1 &*

*prog2 &*

where “&” sends a program to the background.

# Communication between processes

**Processes**\* cannot refer to each others' variables (ensured by the hardware – MMU\*\*), and consequently they do not exhibit the same difficulties as threads.

To achieve concurrency, communication between the programs must occur via mechanisms provided by:

- Operating system via files: *fopen()*, *fread()*, *fwrite()*, *fclose()*;
- Libraries via messages, e.g., *send()*, *get()*, etc.;
- Other mechanisms, e.g., kernel.

\* – **Processes** are imperative programs with their own memory spaces.

\*\* – Memory management unit (protects the memory of one process from accidental reads and writes by another).

# Example of files handling by processes

## #Write to a binary file using fwrite()

```
#include <stdio.h>
#include <stdlib.h>

struct threeNum{
    int n1, n2, n3};

int main(){
    int n;
    struct threeNum num;
    FILE *fptr;

    if ((fptr = fopen("C:\\program.bin", "wb")) == NULL){
        printf("Error! opening file");
        exit(1); // Program exits if the file pointer returns NULL.
    }

    for(n = 1; n < 5; ++n){
        num.n1 = n;
        num.n2 = 5*n;
        num.n3 = 5*n + 1;
        fwrite(&num, sizeof(struct threeNum), 1, fptr);
    }
    fclose(fptr);
    return 0;
}
```

## #Read from a binary file using fread()

```
#include <stdio.h>
#include <stdlib.h>

struct threeNum{
    int n1, n2, n3};

int main(){
    int n;
    struct threeNum num;
    FILE *fptr;

    if ((fptr = fopen("C:\\program.bin", "rb")) == NULL){
        printf("Error! opening file");
        exit(1); // Program exits if the file pointer returns NULL.
    }

    for(n = 1; n < 5; ++n){
        fread(&num, sizeof(struct threeNum), 1, fptr);
        printf("n1: %d\\n2: %d\\n3: %d\\n", num.n1, num.n2, num.n3);
    }
    fclose(fptr);

    return 0;
}
```

# To sum up

- Multitasking is a mid-layer level of concurrency (between hardware and conceptual abstraction) supported by the OS kernel through:
  - Memory allocation;
  - Scheduling of tasks.
- Two main multitasking mechanisms provided by an OS are:
  - Processes – imperative programs with their own memory spaces.
  - Threads – imperative programs that run concurrently and share a memory space.
- Shared memory of threads impose several challenges:
  - Memory consistency – defined from all possible implementation of threads;
  - Race conditions - the order of threads execution affects program results (can be resolved with mutex);
  - Deadlocks – side-effect of mutexes when threads become permanently blocked trying to acquire locks held by other threads;
- Processes are not prone to the above issues until they start communicating with other processes.

# The end!

See you next time – **April 24.**

Well done with the timely submission of Assignment 1!

Late submission for Assignment 1 will be cut-off on **April 20, 23:59.**